

ML-4

1. **Ensemble techniques** combine multiple models to improve predictions (e.g., Random Forest, Boosting).
2. **Bagging** trains models in parallel on bootstrapped samples, then aggregates results (e.g., Random Forest).
3. **Bootstrapping** creates diverse training subsets by sampling with replacement.
4. **Random Forest** uses bagging with decision trees + random feature selection.
5. **Randomization** reduces overfitting by limiting tree correlation (via feature bagging).
6. **Feature bagging** randomly selects subsets of features per tree split.
7. **Decision trees in Gradient Boosting** are weak learners trained sequentially on residuals.
8. **Bagging vs Boosting:**
 - *Bagging*: Parallel, reduces variance.
 - *Boosting*: Sequential, reduces bias.
9. **AdaBoost** adjusts weights for misclassified points in each iteration.
10. **Weak learners** perform slightly better than random guessing (e.g., shallow trees).
11. **Adaptive Boosting** reweights misclassified samples to focus on hard cases.
12. **AdaBoost weight adjustment**: Increases weights for misclassified points.
13. **XGBoost advantages**: Regularization, parallel processing, handling missing data.
14. **XGBoost regularization**: Penalizes complex trees (L1/L2).
15. **Ensemble types**: Bagging, Boosting, Stacking.
16. **Bagging vs Boosting comparison**:
 - *Bagging*: Good for high-variance models.
 - *Boosting*: Good for high-bias models.
17. **Ensemble diversity**: Using varied models/data reduces collective errors.
18. **Improved performance**: Averages errors, reduces overfitting.
19. **Ensemble bias/variance**:
 - *Bias*: Averaging reduces bias (Boosting).
 - *Variance*: Aggregation reduces variance (Bagging).
20. **Trade-off**: More models reduce variance but increase computation.
21. **Applications**: Fraud detection (Random Forest), ranking (XGBoost).

- 22. **Interpretability:** Less interpretable than single models (trade-off for accuracy).
- 23. **Stacking:** Combines models via meta-learner (e.g., logistic regression).
- 24. **Meta-learners:** Train on base model predictions (final aggregator).
- 25. **Challenges:** Computationally expensive, risk of overfitting.
- 26. **Boosting** iteratively corrects errors of prior models (vs bagging's parallelism).
- 27. **Boosting intuition:** Focuses on hard-to-predict samples.
- 28. **Sequential training:** Each model learns from previous errors.
- 29. **Misclassified handling:** Reweights or resamples difficult points.
- 30. **Weights:** Prioritize high-error samples in next iteration.
- 31. **AdaBoost vs Boosting:** AdaBoost is a specific boosting algorithm with weight updates.
- 32. **AdaBoost weight adjustment:** Doubles weights for misclassified points.
- 33. Weak learners: Simple models (e.g., depth-1 trees) slightly better than random.
- 34. Gradient Boosting: Fits trees to residuals (errors) of previous models.
- 35. Gradient descent: Minimizes loss function by iteratively updating model.
- 36. Learning rate: Controls contribution of each tree (smaller = more robust).
- 37. Overfitting handling: Early stopping, shrinkage (learning rate < 0.1).
- 38. XGBoost vs Gradient Boosting: XGBoost adds regularization, parallel processing.
- 39. Regularized boosting: Penalizes tree complexity (L1/L2 in XGBoost).
- 40. XGBoost advantages: Faster, handles missing data, built-in cross-validation.
- 41. Early stopping: Halts training if validation score doesn't improve.
- 42. Prevents overfitting: Stops before model fits noise.

43. Hyperparameters: Learning rate, tree depth, subsampling ratio.
44. Challenges: Sensitive to noise, computationally intensive.
45. Convergence: Stops when errors stop improving significantly.
46. Performance improvement: Sequentially corrects errors.
47. Data imbalance: Can worsen minority class errors (use class weights).
48. Applications: Click prediction (XGBoost), medical diagnosis (AdaBoost).
49. Ensemble selection: Chooses subset of models to optimize performance.
50. Interpretability: Feature importance scores (less transparent than single trees).
51. **Curse of dimensionality**: High dimensions increase sparsity, hurt KNN performance.
52. **KNN applications**: Recommendation systems, image classification.
53. **Weighted KNN**: Closer neighbors have higher voting power.
54. **Missing values**: Impute or use distance metrics handling missingness (e.g., Gower).
55. **Lazy vs eager learning**: KNN is lazy (no training, predicts at runtime).
56. **Improving KNN**: Feature scaling, dimensionality reduction, optimal K .
57. **KNN regression**: Predicts average (or median) of K neighbors.
58. **Decision boundary**: Highly irregular (voronoi tessellation).
59. **Choosing K** : Cross-validation or elbow method (error vs K plot).
60. **Small vs large K** :
 - *Small*: Noisy but flexible.
 - *Large*: Smooth but may underfit.
61. **Feature scaling**: Critical for distance-based metrics (e.g., Min-Max).
62. **KNN vs SVM/DT**:
 - *KNN*: No training, slow prediction.
 - *SVM*: Better for high dimensions.
 - *DT*: More interpretable.
63. **Distance metric impact**: Euclidean (scale-sensitive), Manhattan (robust to outliers).
64. **Imbalanced data**: Use weighted voting or resampling.
65. **Cross-validation**: Tunes K and distance metrics via grid search.
66. **Uniform vs distance-weighted**:
 - *Uniform*: Equal neighbor votes.
 - *Weighted*: Closer neighbors matter more.
67. **Computational complexity**: $O(n)$ per query (slow for large datasets).

68. **Outlier sensitivity:** Manhattan distance is more robust than Euclidean.
69. **Elbow method:** Pick K where error rate stabilizes.
70. **Text classification:** Use TF-IDF vectors + cosine similarity.
71. **PCA components:** Retain 95% variance or use scree plot.
72. **Reconstruction error:** Measures info loss when reducing dimensions.
73. **PCA applications:** Image compression, noise reduction.
74. **PCA limitations:** Linear assumptions, loses interpretability.
75. **SVD:** Generalization of PCA (works on non-square matrices).
76. **LSA:** Applies SVD to term-document matrices (topic modeling).
77. **PCA alternatives:** t-SNE, UMAP, ICA.
78. **t-SNE advantages:** Preserves local structure, good for visualization.
79. **t-SNE vs PCA:** t-SNE captures nonlinear patterns; PCA is linear.
80. **t-SNE limitations:** Computationally heavy, no inverse transform.
81. **PCA vs ICA:** PCA finds uncorrelated components; ICA finds independent sources.
82. **Manifold learning:** Unfolds nonlinear structures (e.g., Isomap).
83. **Autoencoders:** Neural networks for nonlinear dimensionality reduction.
84. **Nonlinear challenges:** Harder to interpret, computationally expensive.
85. **Distance metric impact:** Euclidean for PCA, cosine for text.
86. **Visualization:** Scatter plots (2D/3D), heatmaps.
87. **Feature hashing:** Reduces dimensions via hashing trick (for high-cardinality data).
88. **Global vs local methods:**
 - *Global:* PCA preserves overall structure.
 - *Local:* t-SNE preserves neighborhood relationships.
89. **Sparsity impact:** May require specialized methods (e.g., sparse PCA).
90. **Outlier impact:** Can distort principal components (use robust PCA).